

CSDS341 Project - Airline Querying System - Final Report

Luke Zhang (rxz330), Jiamu Zhang (jxz1217), Quynh Nguyen (qtn2)

2022-07-25

Contents

Background & Introduction	2
Database High-Level Design	2
Assumptions	2
Entity-Relationship Model	3
Data Description	3
Relational Schema	4
Strong Entities	4
Merged Tables	5
Weak Entity	6
Relationships	6
Constraints	7
Primary Key Constraints	7
Overlap Constraints	7
Foreign Key Constraints	8
Mapping Cardinalities	8
Participation Constraints	9
Functional Dependencies	9
Example Queries	11
Preliminaries	11
Set-Up Connection to Database	11
Easy-Level Query	11
Medium-Level Query	12
Hard-Level Query	13

Implementation	14
GUI Interface	15
Roles of members	15
Luke Zhang	15
Jiamu Zhang	15
Quynh Nguyen	15
What we learned from this project	15

Background & Introduction

In recent decades, the growing demand for leisure and business travel has led to the prosperity of the airline market. An increasing number of people have been choosing to take flights to travel domestically or internationally. Therefore, an organized and comprehensive database that stores the airline system is critical for both travelers and crew to obtain plenty of simultaneous information.

Although there do exist several flight databases or applications for commercial airlines, it is rare to find comprehensive information - including weather at the departure airport and destination, aircraft type, the total flight hours of pilots, and the number of luggage allowed - in just one database. This information offers travelers a chance to be better prepared for traveling.

Since our airline querying system contains a relatively extensive data sets, the crew members who choose to use our database are able to access the basic information about the travelers who will be on their flight and provides updates about the airline information.

Database High-Level Design

Assumptions

In order to simulate the real-world situation of a complex flight system, we need to design our database based on the following assumptions. These assumptions can also be used as references for all data constraints.

1. Assume that there have and only have two types of users of the airline querying system: travelers and crew.
2. Assume that plane ticket information is stored in the database system and each ticket is only valid for one traveler. However, a traveler may own zero or more plane tickets.
3. Assume that each ticket contains a specific seat location for exactly one flight. However, a flight may have multiple tickets being sold to travelers.
4. Assume that a crew member can be either an air attendant or a pilot. Therefore, a crew member can serve zero or more flights. Additionally, a flight must be served by at least one crew member. It does have a slight chance that a small propeller airplane only needs one crew member (i.e. the pilot).
5. Assume that a flight is operated by exactly one aeroplane. For example, the aircraft with registration number B-6075 is operating a specific flight (flight number: CA862) from Beijing(PEK) to Geneva(GVA). However, it is likely that one aeroplane can fly multiple flights. Notice that the registration number is unique for each aeroplane.

6. Assume that an aeroplane can only belong to one company. Additionally, a airline company can have multiple planes.
7. Assume that each airline company must have at least one airport as its hub, a place where the head-quarter of the company locates and where the aeroplanes get maintenance and repaired. However, some large airports can provide services for multiple airline companies. For example, Los Angeles International Airport (LAX) is a hub for both United and Delta Airlines, and Delta Airlines has another hub : Detroit Metropolitan Airport (DTW).
8. Assume that each flight can have multiple schedules, and a schedule can be mapped to multiple flights. It is common for most domestic flights to have the same flight flying the same route on two successive days. There is also a slight chance that two flights have the exact same schedule.
9. Assume that each flight only departures from exactly one airport and only arrives at exactly one airport. However, an airport can have many flights.

Entity-Relationship Model

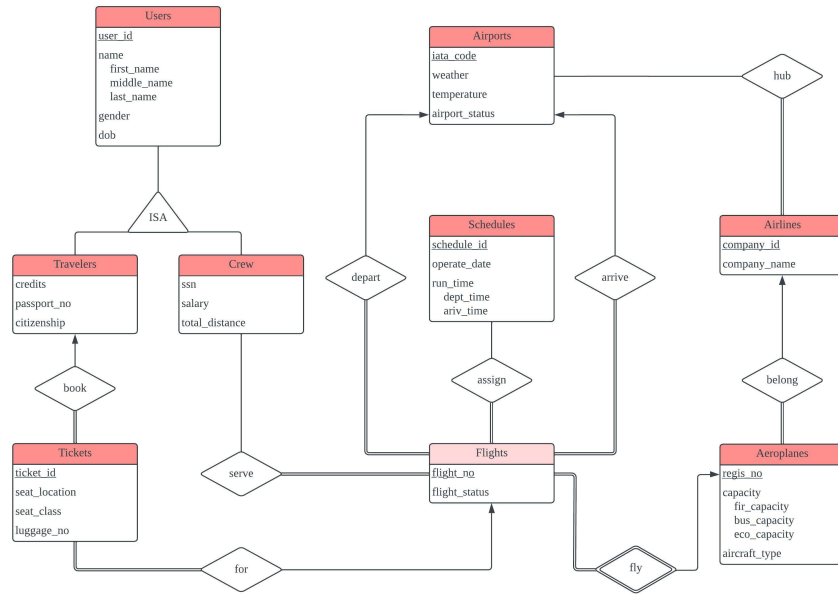


Figure 1: CSDS341 Airline Database ER-Diagram

Data Description

Our database consists of seven entities, eight relationships, and one identifying relationship. Notice that some tables and relationships are merged together to reduce redundancy.

Travelers and **Crew** are considered as users of our airline user system. **Travelers** contains each travelers' information such as **credits**, **citizenship**, etc. **Crew**, sharing part of the attributes with **Travelers**, also has its unique attributes such as **ssn** and **salary**. Note that each user should be either a Traveler or a Crew. A user cannot be both or none, which means that if a **user_id** has already been stored in **Travelers** table, it will not be present in the **Crew** Table.

For **Travelers**, they are connected to their booking information (i.e. Table **Tickets**) by the **book** relationship, while travelers connect to **flight** information by the **serve** relationship. Note that **flight** cannot

exist by itself, which means that a flight must be operated by an aeroplane. Additionally, an aeroplane must belong to exactly one airline company by the **belong** relationship. Moreover, a flight is connected to two airports by the **arrive** and **depart** relationships. An airline company can also connect to **airports** by the **hub** relationship.

Our database implement NOT NULL constraints for primary keys and part of other attributes. We also use FOREIGN KEY constraints when referencing information from other tables. PRIMARY KEY constraints are adopted in all tables to ensure that each tuple is unique. PRIMARY KEY constraints also prevent the NULL values being assigned into the table by default. For more information, please see the constraints section where our group listed all primary key constraints, overlap Constraints, Foreign Key Constraints, Mapping Cardinalities, Participation Constraints.

For Functional Dependencies and Normalization, all of our tables are in BCNF. For each table, the functional dependencies are non-trivial and for each FD $X \rightarrow Y$, X is one of the super key of the table. Since there are no violation of BCNF, no tables should be decomposed from our current design. We will discussed our functional dependencies in more detail later in the Functional Dependencies Section.

Relational Schema

Strong Entities

The **Travelers** Table records the information of each traveler such as citizenship and passport numbers.

```
CREATE TABLE IF NOT EXISTS Travelers (  
    user_id INT NOT NULL,  
    first_name VARCHAR(50),  
    middle_name VARCHAR(50),  
    last_name VARCHAR(50),  
    gender CHAR(1),  
    dob DATE,  
    credits INT DEFAULT 0,  
    passport_no VARCHAR(20),  
    citizenship VARCHAR(30),  
    PRIMARY KEY (user_id),  
    CHECK (gender IN ('M' , 'F', 'U')));
```

The **Crew** Table is built for crew members including pilots. Their information such as social-security number and total distance traveled will be recorded.

```
CREATE TABLE IF NOT EXISTS Crew (  
    user_id INT NOT NULL,  
    first_name VARCHAR(50),  
    middle_name VARCHAR(50),  
    last_name VARCHAR(50),  
    gender CHAR(1),  
    dob DATE,  
    ssn INT NOT NULL,  
    salary DOUBLE,  
    total_distance INT,  
    PRIMARY KEY (user_id),  
    CHECK (gender IN ('M' , 'F', 'U')));
```

The **Airports** Table contains information of airports such as the IATA code and name of the airport.

```
CREATE TABLE IF NOT EXISTS Airports (
    iata_code CHAR(3) NOT NULL,
    airport_name VARCHAR(100),
    country CHAR(2),
    weather VARCHAR(30),
    airport_status VARCHAR(30),
    PRIMARY KEY (iata_code),
    CHECK (weather IN ('Sunny' , 'Mostly Sunny', 'Partly Cloudy',
        'Cloudy', 'Rainy', 'Heavy Rainy', 'Foggy',
        'Snowy', 'Heavy Snowy', 'Frost')),
    CHECK (airport_status IN ('Free', 'Normal',
        'Busy', 'Small-Scale Delay', 'Large-Scale Delay')));
```

The Airlines Table contains information of airline companies.

```
CREATE TABLE IF NOT EXISTS Airlines (
    company_id INT NOT NULL,
    company_name VARCHAR(50),
    PRIMARY KEY (company_id));
```

The Schedules Table contains information of flight schedules including departure time and arrival time.

```
CREATE TABLE IF NOT EXISTS Schedules (
    schedule_id INT NOT NULL,
    dept_date DATE,
    dept_time TIME,
    ariv_date DATE,
    ariv_time TIME,
    PRIMARY KEY (schedule_id));
```

Merged Tables

The Aeroplanes_belong Table merges one entity (Aeroplanes) and one relationship (belong). It contains information about connection between aeroplanes and airline companies.

```
CREATE TABLE IF NOT EXISTS Aeroplanes_belong (
    regis_no VARCHAR(10) NOT NULL,
    fir_capacity INT,
    bus_capacity INT,
    eco_capacity INT,
    aircraft_type VARCHAR(50),
    company_id INT NOT NULL,
    PRIMARY KEY (regis_no),
    FOREIGN KEY (company_id)
        REFERENCES Airlines(company_id));
```

The Tickets_book_for Table merges one entity (Tickets) and two relationships (book and for). It contains traveler booking information.

```
CREATE TABLE IF NOT EXISTS Tickets_book_for (
    ticket_id INT NOT NULL,
```

```

seat_location CHAR(4),
seat_class CHAR(1),
luggage_no INT,
regis_no VARCHAR(10) NOT NULL,
flight_no VARCHAR(7) NOT NULL,
traveler_id INT NOT NULL,
PRIMARY KEY (ticket_id),
FOREIGN KEY (traveler_id)
    REFERENCES Travelers(user_id),
FOREIGN KEY (regis_no, flight_no)
    REFERENCES Flights_ariv_dept(regis_no, flight_no),
CHECK (seat_class IN ('F' , 'B', 'E')));

```

Weak Entity

The `Flights_ariv_dept` Table merges one weak entity (`Flights`), one identifying relationship (`Fly`), and two relationships (`arrive`, `depart`). It contains the information of flights as well as its departure and arrival information.

```

CREATE TABLE IF NOT EXISTS Flights_ariv_dept (
    regis_no VARCHAR(10) NOT NULL,
    flight_no VARCHAR(7) NOT NULL,
    flight_status VARCHAR(10),
    dept_iata_code CHAR(3) NOT NULL,
    ariv_iata_code CHAR(3) NOT NULL,
    PRIMARY KEY (regis_no, flight_no),
    FOREIGN KEY (regis_no)
        REFERENCES Aeroplanes_belong(regis_no),
    FOREIGN KEY (dept_iata_code)
        REFERENCES Airports(iata_code),
    FOREIGN KEY (dept_iata_code)
        REFERENCES Airports(iata_code),
    CHECK (flight_status IN ('On-Time' , 'Delay', 'Cancel')));

```

Relationships

The `serve` Table contains information of connection between Crew members and the flights they serve.

```

CREATE TABLE IF NOT EXISTS serve (
    crew_id INT NOT NULL,
    regis_no VARCHAR(10) NOT NULL,
    flight_no VARCHAR(7) NOT NULL,
    PRIMARY KEY (crew_id , regis_no , flight_no),
    FOREIGN KEY (crew_id)
        REFERENCES Crew(user_id),
    FOREIGN KEY (regis_no, flight_no)
        REFERENCES Flights_ariv_dept(regis_no, flight_no));

```

The `assign` Table contains information of connection between flights and flying schedules.

```

CREATE TABLE IF NOT EXISTS assign (

```

```

    regis_no VARCHAR(10) NOT NULL,
    flight_no VARCHAR(7) NOT NULL,
    schedule_id INT NOT NULL,
    PRIMARY KEY (regis_no , flight_no , schedule_id),
    FOREIGN KEY (regis_no, flight_no)
        REFERENCES Flights_ariv_dept(regis_no, flight_no),
    FOREIGN KEY (schedule_id)
        REFERENCES Schedules(schedule_id));

```

The hub Table contains information of connection between airline companies and their hub airports.

```

CREATE TABLE IF NOT EXISTS hub (
    company_id INT NOT NULL,
    iata_code CHAR(3) NOT NULL,
    PRIMARY KEY (company_id , iata_code),
    FOREIGN KEY (company_id)
        REFERENCES Airlines(company_id),
    FOREIGN KEY (iata_code)
        REFERENCES Airports(iata_code));

```

Constraints

Primary Key Constraints

The Primary Key of each table can either be one or several attributes. Assume that a primary key is NOT NULL and unique.

1. For Table **Travelers** and **Airports**, the primary key is **user_id**. Since a traveler or a crew **isa** user, a **user_id** cannot be used for a traveler or a crew simultaneously, which will also be discussed in the Overlap Constraints Section.
2. For Table **Airports**, the primary key is **iata_code**. It stands for the unique three-character International Air Transport Association Code for each airport.
3. For Table **Airlines**, the primary key is **company_id**.
4. For Table **Schedules**, the primary key is **schedule_id**.
5. For Table **Aeroplanes_belong**, the primary key is **regis_no**.
6. For Table **Flights_ariv_dept**, the primary keys are **regis_no** and **flight_no**.
7. For Table **Tickets_book_for**, the primary key is **ticket_id**.
8. For Table **serve**, the primary keys are **crew_id**, **regin_no** and **flight_no**.
9. For Table **assign**, the primary keys are **regis_no**, **flight_no** and **schedule_id**.
10. For Table **hub**, The primary keys are **company_id** and **iata_code**.

Overlap Constraints

A **User** can be either **Travelers** or **Crew**, but they share the same primary key attribute **user_id**, so a **user_id** cannot be used for a traveler or a crew simultaneously.

Foreign Key Constraints

1. The `Tickets_book_for` relation has three foreign-key constraints:
 - `traveler_id` references the primary key of the `Travelers` relation (`Travelers.user_id`).
 - `regis_no` references the primary key of the `Flights_ariv_dept` relation (`Flights_ariv_dept.regis_no`).
 - `flight_no` reference the primary key of the `Flights_ariv_dept` relation (`Flights_ariv_dept.flight_no`).
2. The `Aeroplanes_belong` relation has one foreign-key constraint:
 - `company_id` references the primary key of the `Airlines` relation (`Airlines.company_id`).
3. The `Flights_ariv_dept` relation has three foreign-key constraints:
 - `regis_no` references the primary key of the `Aeroplanes_belong` relation (`Aeroplanes_belong.regis_no`).
 - `dept_iata_code` references the primary key of the `Airports` relation (`Airports.iata_code`).
 - `ariv_iata_code` references the primary key of the `Airports` relation (`Airports.iata_code`).
4. The `serve` relation has three foreign-key constraints:
 - `crew_id` references the primary key of the `Crew` relation (`Crew.user_id`).
 - `regis_no` references the primary key of the `Flights_ariv_dept` relation (`Flights_ariv_dept.regis_no`).
 - `flight_no` reference the primary key of the `Flights_ariv_dept` relation (`Flights_ariv_dept.flight_no`).
5. The `assign` relation has three foreign-key constraints:
 - `schedule_id` references the primary key of the `Schedules` relation (`Schedules.schedule_id`).
 - `regis_no` references the primary key of the `Flights_ariv_dept` relation (`Flights_ariv_dept.regis_no`).
 - `flight_no` reference the primary key of the `Flights_ariv_dept` relation (`Flights_ariv_dept.flight_no`).
6. The `hub` relation has three foreign-key constraints:
 - `company_id` references the primary key of the `Airlines` relation (`Airlines.company_id`).
 - `iata_code` references the primary key of the `Airports` relation (`Airports.iata_code`).

Mapping Cardinalities

1. `book` (merged in relation `Tickets_book_for`) is a many-to-one relationship set from `Tickets` to `Travelers`, meaning that a travelers can book multiple tickets for flights and a ticket can only be used for one traveler.
2. `for` (merged in relation `Tickets_book_for`) is a many-to-one relationship set from `Tickets` to `Flights`, meaning that a flight can have many tickets being sold and one ticket can only be used for boarding one flight.
3. `serve` is a many-to-many relationship set from `Crew` to `Flights`, meaning that several crews can work on the same flight, and one crew can work on multiple flights.
4. `depart` and `arrive` (merged in relation `Flights_ariv_dept`) are many-to-one relationships from `Flights` to `Airports`, meaning that one airport can have many departing or arriving aeroplanes and one flight can only depart (or arrive) at one airport.
5. `assign` is a many-to-many relationship, meaning that a flight can have several schedules. For instance, a flight can be operated on two successive days. Additionally, for one schedule, there might be two or more flights being operated.
6. `belong` is a many-to-one relationship set from `Aeroplanes` to `Airlines`, meaning that one airline company can own multiple planes, while a plane can only belong to one airline company.
7. `hub` is a many-to-many relationship set from `Airlines` to `Airports`, meaning that airline companies can have several hub airports, and one airport can be the hub airport of multiple airlines companies.

Participation Constraints

1. **Tickets** total participates in **book** and **for**. A ticket must be assigned to one traveler and must be mapped to a flight. However, a traveler can have no tickets and a flight can have no tickets being sold.
2. **Flights** total participates in the **depart**, **arrive** and **assign**. This means that a ticket must be assigned to a certain schedule as well as a departure airport and an arrival airport. However, an airport can have no departing or arriving airports, and a schedule can have no flights in that schedule.
3. **Flights** total participates in **serve**. This means that a flight must be served by at least one crew member, while a crew member can serve zero flights.
4. **Aeroplanes** total participates in **belong**. This means that an aeroplane must be owned by one airline company and one airline company can own zero aeroplanes.
5. **Airlines** total participates in **hub**. This means that an airline must have at least one airport as its hub and one airport can have zero or many airlines.

Functional Dependencies

Travelers Table

$$F_{\text{travelers}} = \{\text{user_id} \rightarrow \{\text{all other attributes}\}, \{\text{passport_no}, \text{citizenship}\} \rightarrow \{\text{all other attributes}\}\}$$

There are two functional dependencies for **travelers**. **user_id** and **{passport_no, citizenship}** are two super keys of the **travelers** relation, which means that the functional dependencies do not violate the BCNF. Note that there is only one candidate key which is **user_id** because **{passport_no, citizenship}** is not minimal.

Crew Table

$$F_{\text{crew}} = \{\text{user_id} \rightarrow \{\text{all other attributes}\}, \text{ssn} \rightarrow \{\text{all other attributes}\}\}$$

There are two functional dependencies for **travelers**. **user_id** and **ssn** are two super keys of the **travelers** relation, which means that the functional dependencies do not violate the BCNF. Note that both **user_id** and **ssn** are candidate keys.

Airports Table

$$F_{\text{airports}} = \{\text{iata_code} \rightarrow \{\text{all other attributes}\}, \{\text{airport_name}, \text{country}\} \rightarrow \{\text{all other attributes}\}\}$$

There are two functional dependencies for **airports**. **iata_id** and **{airport_name, country}** are two super keys of the **airports** relation, which means that the functional dependencies do not violate the BCNF.

Airlines Table

$$F_{\text{airlines}} = \{\text{compnay_id} \rightarrow \{\text{company_name}\}\}$$

There is just one functional dependency for **airports**. **iata_id** is the super key of the **airports** relation, which means that the functional dependency does not violate the BCNF.

Schedules Table

$$F_{\text{Schedules}} = \{\text{schedule_id} \rightarrow \{\text{all other attributes}\}\}$$

There is just one functional dependency for **Schedules**. `schedule_id` is the super key of the **Schedules** relation, which means that the functional dependency does not violate the BCNF.

Aeroplanes_belong Table

$$F_{\text{Aeroplanes_belong}} = \{\text{regis_no} \rightarrow \{\text{all other attributes}\}\}$$

There is just one functional dependency for **Aeroplanes_belong**. `regis_no` is the super key of the **Aeroplanes_belong** relation, which means that the functional dependency does not violate the BCNF.

Tickets_book_for Table

$$F_{\text{Tickets_book_for}} = \{\text{ticket_id} \rightarrow \{\text{all other attributes}\}\}$$

There is just one functional dependency for **Tickets_book_for**. `ticket_id` is the super key of the **Tickets_book_for** relation, which means that the functional dependency does not violate the BCNF.

Flights_ariv_dept Table

$$F_{\text{Flights_ariv_dept}} = \{\{\text{regis_no}, \text{flight_no}\} \rightarrow \{\text{all other attributes}\}\}$$

There is just one functional dependency for **Flights_ariv_dept**. `ticket_id` is the super key of the **Flights_ariv_dept** relation, which means that the functional dependency does not violate the BCNF.

serve Table

$$F_{\text{serve}} = \{\{\text{regis_no}, \text{flight_no}, \text{crew_id}\} \rightarrow \{\text{all attributes}\}\}$$

There is just one functional dependency for **serve**. `{regis_no, flight_no, crew_id}` is the super key of the **serve** relation, which means that the functional dependency does not violate the BCNF.

assign Table

$$F_{\text{assign}} = \{\{\text{regis_no}, \text{flight_no}, \text{schedule_id}\} \rightarrow \{\text{all attributes}\}\}$$

There is just one functional dependency for **assign**. `{regis_no, flight_no, crew_id}` is the super key of the **assign** relation, which means that the functional dependency does not violate the BCNF.

hub Table

$$F_{\text{hub}} = \{\{\text{iata_code}, \text{company_id}\} \rightarrow \{\text{all attributes}\}\}$$

There is just one functional dependency for **hub**. `{iata_code, company_id}` is the super key of the **hub** relation, which means that the functional dependency does not violate the BCNF.

Example Queries

Preliminaries

Database connection and data querying will require the following packages.

```
library(RMySQL)
library(RODBC)
library(knitr)
library(tidyverse)
```

Set-Up Connection to Database

The following code set-up a connection to our airline database built by MySQL. The function `dbListTables` function lists the name of all tables we have in the database alphabetically.

```
# Establish Connection to MySQL
sql_link <- dbConnect(RMySQL::MySQL(),
                      dbname='airline',
                      host='localhost',
                      port=3306,
                      user='root',
                      password='Lukezrh@010408')

# List All the Tables in the Schema
dbListTables(sql_link)
```

```
## [1] "aeroplanes_belong" "airlines"          "airports"
## [4] "assign"            "crew"              "flights_ariv_dept"
## [7] "hub"               "schedules"         "serve"
## [10] "tickets_book_for"  "travelers"
```

Easy-Level Query

Query 01: Find the first name(`first_name`), last name(`last_name`), date of birth(`dob`), and the credits(`credits`) of all male travelers who are U.S. citizen.

SQL & Output Table:

We convert the Query to SQL and Stored in the variable `sql_01` and generate the following table as the output.

```
sql_01 <- "SELECT first_name, last_name, dob, credits
          FROM   travelers
          WHERE  citizenship = 'US' AND
                gender = 'M' AND
                dob < '1985-01-01'"

result_01 <- dbSendQuery(sql_link, sql_01)
output_01 <- fetch(result_01)
output_01 %>% kable()
```

first_name	last_name	dob	credits
Adrian	Rosales	1982-05-05	6269
Zachary	Shipley	1982-09-17	1776
Sean	Kim	1981-12-07	4915
Jose	Villatoro	1981-11-13	1645
Daniel	Ramirez	1968-10-14	6937
Turki	al-Ahsan	1965-01-01	8458
Andrew	Macy	1968-09-20	9060
Muhaajir	el-Noorani	1949-04-04	5074
Sterling	Kirkland	1961-03-06	4279
Darius	Stukes	1952-03-30	9913
Deven	Granderson	1951-04-25	5676

Relational Algebra:

$$\begin{aligned} \text{usmale} &\leftarrow \sigma_{\text{gender}='M' \wedge \text{citizenship}='US'}(\text{travelers}) \\ \text{result} &\leftarrow \Pi_{\text{first_name}, \text{last_name}, \text{dob}, \text{credits}}(\sigma_{\text{dob} < '1985-01-01'}(\text{usmale})) \end{aligned}$$

Tuple Relational Calculus:

$$\{t^4 \mid (\exists p)(p \in \text{travelers} \wedge t[1] = p[\text{first_name}] \wedge t[2] = p[\text{last_name}] \wedge t[3] = p[\text{dob}] \wedge t[4] = p[\text{credits}] \wedge t[\text{citizenship}] = 'us' \wedge t[\text{gender}] = 'M' \wedge t[\text{dob}] < '1985-01-01')\}$$

Medium-Level Query

Query 02: Find the flight number(flight_no), departure airport(dept_iata_code), arrival airport(ariv_iata_code) of all flights that operated by China Eastern Airlines and has a capacity of 30 seats in first class cabin.

SQL & Output Table:

```
sql_02 <- "SELECT flight_no, dept_iata_code, ariv_iata_code
FROM Flights_ariv_dept NATURAL JOIN Aeroplanes_belong
WHERE (Aeroplanes_belong.fir_capacity = 30)
AND (regis_no IN (
SELECT regis_no
FROM Aeroplanes_belong NATURAL JOIN Airlines
WHERE Airlines.company_name = 'China Eastern Airlines'))"

result_02 <- dbSendQuery(sql_link, sql_02)

output_02 <- fetch(result_02)
output_02 %>% kable()
```

flight_no	dept_iata_code	ariv_iata_code
MU1332	ORF	SJC
MU3212	VBG	KEF
MU3762	GSP	CEB
MU470	GSP	ISE

flight_no	dept_iata_code	ariv_iata_code
MU5879	NNG	BDL
MU6141	DAY	RND
MU7451	FOC	AFW
MU8560	KWL	CAE
MU8990	PHL	SDQ
MU9025	HSV	OTP
MU9871	CAE	ITM
MU2680	ORF	DJJ
MU4422	SHE	AYT
MU4868	VBG	ADA
MU4936	DAY	SJC
MU636	LIT	CRK
MU7946	TYN	ERI
MU8048	RDU	TSE
MU8188	GSP	BLA
MU8497	LIT	MCI
MU8723	MBS	DCA
MU9773	PHL	TPA

Relational Algebra:

Tuple Relational Calculus:

Hard-Level Query

Query 03: Find the passport number(passport_no) and first name(first_name) of all Chinese Traveler who has already booked at least one Plane Ticket. The result table should also include citizenship as a check.

SQL & Output Table:

```
sql_03 <- "SELECT passport_no, first_name, citizenship
          FROM   Travelers
          WHERE  (Citizenship = 'CN')
          AND
          (user_id IN (SELECT   traveler_id
                        FROM     Tickets_book_for
                        GROUP BY  traveler_id
                        HAVING    count(ticket_id) >= 1))"

result_03 <- dbSendQuery(sql_link, sql_03)

output_03 <- fetch(result_03)
output_03 %>% kable()
```

passport_no	first_name	citizenship
UU836955458	Jesus	CN
IB430482154	Brandon	CN
OQ380375182	Haseena	CN
LM812170185	John	CN

passport_no	first_name	citizenship
EB949426590	Paul	CN
TL625875287	Jordan	CN
NJ572174743	Natelle Rae	CN
TT276365122	Brandon	CN
GA510923640	Fatima	CN
RX309492405	Sidney	CN
KE075171969	Kristina	CN
BH026711943	Keisha	CN
CB965376246	Haley	CN
OC970861821	Mariah	CN
IC641671105	Lauren	CN
FX285795793	Daisy	CN
VB903190894	Sierra	CN
BQ269116750	Timothy	CN
FM884378913	Paris	CN
LV473086926	Fat'hi	CN
DQ891369571	Shajee'a	CN
PI055048751	Macennia	CN
LM354095793	Shaakir	CN
BX613479733	Adham	CN

Relational Algebra:

Tuple Relational Calculus:

Implementation

Our database system was implemented as a simple command-line interface that can give an expected output table of queries. We use **MySQL Benchwork 8.0 CE** to create our database. There are eleven tables in our database. The SQL language that we used to create them are documented previously in the schema section.

For the data used in our database system, one of our largest table **Tickets_book_for** contains over 10,000 tuples, while one of our smallest table **Airlines** contains only approx. 50 tuples. Some data are obtained directly from internet. For instance, the airport names and airport IATA codes are obtained from internet, but filtered using **R** programming language to ensure the data quality. However, other part of data are not accessable from internet since they are protected by law or are part of the commercial secrets of airline companies. We use **R** to randomly simulate this type of data which includes the departure time of flights, flight numbers, travelers' passport number, and airport weather data.

For connecting the database from outside, we decide to use **R** programming language. We use mainly **RMySQL** and **RODBC** packages to push our query to **MySQL** server, and to generate the output table as dataframe in **R** environment.

GUI Interface

Roles of members

Luke Zhang

ER Diagram, Constraints, Relational Schema, Functional Dependencies, Example Query (SQL), Implementing R-MySQL Connection

Jiamu Zhang

Constraints, Example Query (TRC, RA), Implementing Java GUI

Quynh Nguyen

Background, Prepare Data for Database, Relational Schema

What we learned from this project